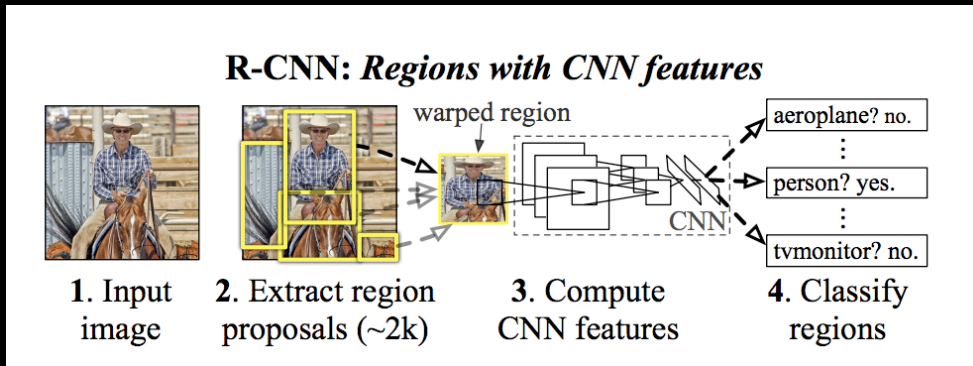


R-CNN (Regions with Convolutional Neural Networks) is a deep learning-based object detection framework introduced by Ross Girshick et al. in 2014. It was one of the first approaches to successfully apply Convolutional Neural Networks (CNNs) for object detection.



Old Object Detection Techniques

- 1) HOG
 - 2) SIFT
 - 3) Selective Search
- * not very accurate or efficient

Localization + Classification

- 1) Region Proposal Methods (Selective Search)
 - ↳ generate potential candidate object locations
- 2) CNN Feature Extraction

Pipeline

1) Region Proposal Generation

- Selective Search 2K

2) Feature Extraction by CNN

- VGG / Alexnet

Resize $(224 \times 224 \times 3)$

3) Classification

- SVM

4) Bounding Box Regression

- Fine Tune Bbox coordinates

Limitations

1) Slow Inference

* Computationally expensive

2000 Regions

50 sec per image on GPU

2) Storage Overhead

* High Storage Requirements

3) Not End to End Trainable

3 components (Region Proposal, Feature
Extraction, Classification)
Trained Separately

Improvements

1) FAST RCNN

2) FASTER RCNN